



## Lab: Terraform State Backend Authentication

The `local` backend stores state as a local file on disk, but other backend types store state in a remote service of some kind, which allows multiple people to access it. Accessing state in a remote service generally requires some kind of access credentials since state data contains extremely sensitive information. It is important to strictly control who can access your Terraform backend.

We will look at two different backend types compatible with Terraform and how each handles authentication.

- Task 1: Authentication: S3 Standard Backend
- Task 2: Authentication: Remote Enhanced Backend

### Backend Configuration: Authentication

Some backends allow us to provide access credentials directly as part of the configuration. However, in normal use we do not recommend including access credentials as part of the backend configuration. Instead, leave those arguments completely unset and provide credentials via the credentials files or environment variables that are conventional for the target system, as described in the documentation for each backend.

### Task 1: Authentication: S3 Standard Backend

The terraform backend end configuration for a given working directory is specified in the Terraform configuration block. Our terraform configuration block for this lab is located in the `terraform.tf` file.

The `s3` backend stores Terraform state as a given key in a given bucket on Amazon S3. This backend supports several different methods in which the Terraform CLI can authenticate to the Amazon S3 bucket.

#### Step 1.1 - Create the bucket in AWS

Create a new bucket within AWS to centrally store your terraform state files:





Amazon S3 > Create bucket

## Create bucket [Info](#)

Buckets are containers for data stored in S3. [Learn more](#)

### General configuration

Bucket name

Bucket name must be unique and must not contain spaces or uppercase letters. [See rules for bucket naming](#)

AWS Region

US East (N. Virginia) us-east-1

Copy settings from existing bucket - *optional*  
Only the bucket settings in the following configuration are copied.

[Choose bucket](#)

### Block Public Access settings for this bucket

Public access is granted to buckets and objects through access control lists (ACLs), bucket policies, access point policies, or all. In order to ensure that public access to this bucket and its objects is blocked, turn on Block all public access. These settings apply only to this bucket and its access points. AWS recommends that you turn on Block all public access, but before applying any of these settings, ensure that your applications will work correctly without public access. If you require some level of public access to this bucket or objects within, you can customize the individual settings below to suit your specific storage use cases. [Learn more](#)

☒ **Block all public access**  
Turning this setting on is the same as turning on all four settings below. Each of the following settings are independent of one another.

☒ **Block public access to buckets and objects granted through *new* access control lists (ACLs)**

**Figure 1:** S3 Remote State Storage

If you're using S3 as a backend, you'll want to configure an IAM policy that solely grants access to the S3 bucket for production to a small handful of trusted people or perhaps solely just the CI server you use to deploy to your environments.

### Step 1.2 - Remove existing resources with terraform destroy

If you already have a state file present with infrastructure deployed from previous labs we will first issue a cleanup of our infrastructure using a `terraform destroy` before changing our our backend.





This is not a requirement as Terraform supports the migration of state data between backends, which will be covered in a future lab.

```
terraform destroy
```

Do you really want to destroy all resources?

Terraform will destroy all your managed infrastructure, as shown above.  
There is no undo. Only 'yes' will be accepted to confirm.

Enter a value: yes

### Step 1.2 - Update Terraform Configuration to use s3 backend

Update the `terraform.tf` configuration to utilize the `s3` backend with the required arguments.

```
terraform {  
  backend "s3" {  
    bucket = "myterraformstate"  
    key    = "path/to/my/key"  
    region = "us-east-1"  
  }  
}
```

Example:

```
terraform {  
  backend "s3" {  
    bucket = "my-terraform-state-ghm"  
    key    = "prod/aws_infra"  
    region = "us-east-1"  
  }  
}
```

Note: A Terraform configuration can only specify a single backend. If a backend is already configured be sure to replace it. Copy just the `backend` block above and not the full `terraform` block  
You can validate the syntax is correct by issuing a `terraform validate`

### Step 1.3 - Provide Terraform AWS credentials to connect to S3 Bucket

Providing credentials for accessing state in an S3 bucket can be done in a number of different ways. This lab will showcase using environment variables but a shared credentials file can also be used. It is important to protect any credential information so while it is possible to set these values in the code itself it is strongly not recommended.





Source credentials to the S3 backend using the `AWS_ACCESS_KEY_ID` and `AWS_SECRET_ACCESS_KEY` environment variables.

```
export AWS_ACCESS_KEY_ID="YOUR_AWS_ACCESS_KEY_ID"
export AWS_SECRET_ACCESS_KEY="YOUR_AWS_SECRET_ACCESS_KEY"
```

(Optional): If using Multi-Factor Authentication you can source the token using the `AWS_SESSION_TOKEN` environment variable.

(Optional): In lieu of setting environment variables you can also utilize an AWS shared credentials file by specifying the path to the file using the `shared_credentials_file` argument within the backend configuration block. This defaults to `~/.aws/credentials`.

#### Step 1.4 - Verify Authentication to S3 Backend

Once the configuration is complete, you can verify authentication to the S3 backend by first removing infrastructure that has already been deployed with a `terraform destroy` and performing a `terraform init`

```
terraform init

Initializing the backend...

Successfully configured the backend "s3"! Terraform will automatically
use this backend unless the backend configuration changes.
```

#### Step 1.5 - Write Terraform State to S3 Backend

Now that authentication has been verified, we can build out the infrastructure with our S3 backend for storing state we can issue a `terraform apply`

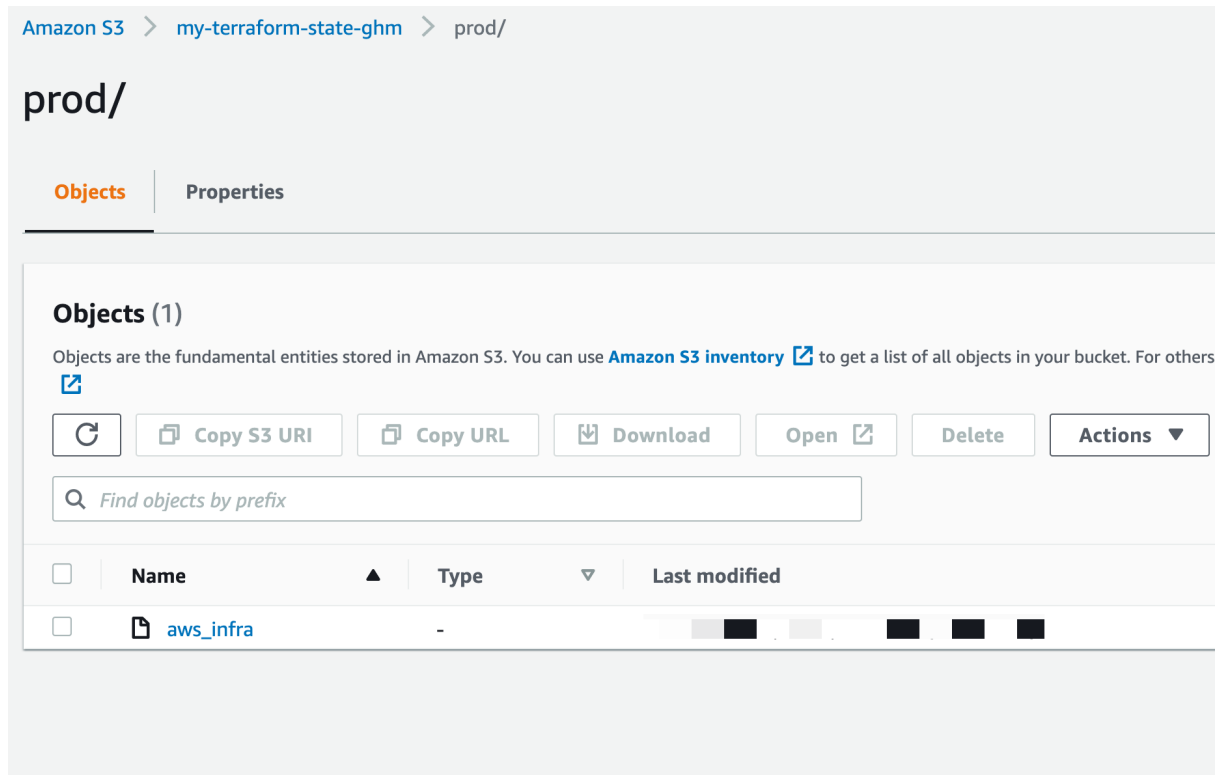
```
terraform apply

...

Do you want to perform these actions?
  Terraform will perform the actions described above.
  Only 'yes' will be accepted to approve.

Enter a value: yes
```





**Figure 2:** S3 Backend Object

Now the state file for your infrastructure build out is stored remotely on the S3 object store in your bucket. This can now be utilized by others who have appropriate permissions to the S3 bucket as we have successfully centralized the terraform state file.

### Step 1.6 - Remove Terraform AWS credentials to connect to S3 Bucket

If you would like to remove access to a centralized state file, you can modify the credentials to your S3 bucket. To showcase unauthenticated access, let's change the source credentials to the S3 backend by incorrectly setting the `AWS_ACCESS_KEY_ID` and `AWS_SECRET_ACCESS_KEY` environment variables.

```
export AWS_ACCESS_KEY_ID="notvalid"
export AWS_SECRET_ACCESS_KEY="notvalid"
```

```
terraform state list
```

```
Error: error configuring S3 Backend: error validating provider credentials: error calling
```





```
status code: 403, request id: 00771614-5553-4eaf-ae8d-a3ce54c66060
```

Once we change these back, you can see we again have access.

```
export AWS_ACCESS_KEY_ID="YOUR_AWS_ACCESS_KEY_ID"  
export AWS_SECRET_ACCESS_KEY="YOUR_AWS_SECRET_ACCESS_KEY"
```

```
terraform state list
```

## Task 2: Authentication: Remote Enhanced Backend

The Terraform `remote` backend stores Terraform state and may be used to run operations in Terraform Cloud. This backend supports the ability to store Terraform state information and perform operations all within Terraform Cloud, based on privileged access.

### Step 2.1

For this task you will have to sign up for a Terraform Cloud account. In order to store state remotely on Terraform Cloud using the `remote` backend we need to create a user token and configure our local environment to utilize that token. We will leverage the `terraform login` command to obtain and save an API token for authenticating to Terraform Cloud.

### Step 2.2

**Note:** You can skip this step if you've already created a organization.

Log in to Terraform Cloud and go to the new organization page:

- New users are automatically taken to the new organization page.
- If your user account is already a member of an organization, open the organization switcher menu in the top navigation bar and click the "Create new organization" button.

Enter a unique organization name and an email address for notifications, then click the "Create organization" button.





## Create a new organization

Organizations are privately shared spaces for teams to collaborate on infrastructure. [Learn more](#) about organizations in Terraform Cloud.

### Organization name

e.g. company-name

Organization names must be unique and will be part of your resource names used in various tools, for example `hashicorp/www-prod`.

### Email address

gabe@maentz.net

The organization email is used for any future notifications, such as billing alerts, and the organization avatar, via [gravatar.com](#).

Create organization

**Figure 3:** New Organization

### Step 2.3

Terraform's CLI needs credentials before it can access Terraform Cloud. We will leverage the `terraform login` command to perform our login to TFC.

```
terraform login
```

Terraform will request an API token **for** `app.terraform.io` using your browser.

If **login** is successful, Terraform will store the token **in** plain text **in** the following file **for** use by subsequent commands:  
`/Users/gabe/.terraform.d/credentials.tfrc.json`

Do you want to proceed?

Only **'yes'** will be accepted to confirm.

Enter a value: yes





Open the following URL to access the tokens page **for** app.terraform.io:  
`https://app.terraform.io/app/settings/tokens?source=terraform-login`

Generate a token using your browser, and copy-paste it into this prompt.

Terraform will store the token **in** plain text **in** the following file  
**for** use by subsequent commands:  
`/home/nyl/.terraform.d/credentials.tfrc.json`

Token **for** app.terraform.io:  
Enter a value:


Retrieved token **for** user gabe\_maentz

Success! Terraform has obtained and saved an API token.

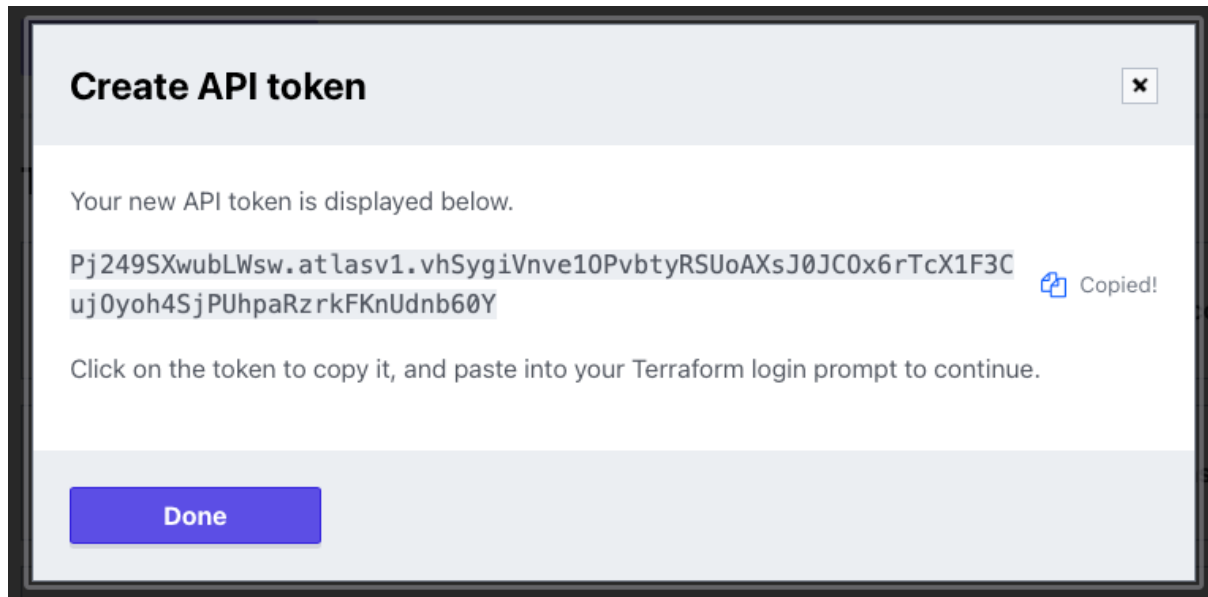
The new API token will be used **for** any future Terraform **command** that must make authenticated requests to app.terraform.io.

**Figure 4:** TFC\_TOKEN

Generate a token using your browser, and copy-paste it into this prompt.







**Figure 5:** TFC\_TOKEN





-----  
-----  
-  
  
New to TFC? Follow these steps to instantly apply an example configuration:

```
$ git clone https://github.com/hashicorp/tfc-getting-started.git  
$ cd tfc-getting-started  
$ scripts/setup.sh
```

We will leverage this `remote` backend authentication to interact with Terraform Cloud from the CLI of our working directory in future labs.

